

RPI ns-2 Graphing and Statistics Package

User Manual and Tutorial

David Harrison
RPI CS Department*

April 29, 2007

Abstract

This user manual contains an installation guide and a short tutorial for the RPI graphing tools for ns.

1 Introduction

The RPI ns-2 Graphing and Statistics Package¹ provides a set of classes for generating commonly used plots and gathering commonly important statistics. The graph package abstracts data collection from rendering. Each graph object instantiates a set of data collection objects which are attached to components in a network topology, as the simulation runs the data collection objects collect data often outputting the data to an intermediate file. After simulation, the graph object conforms the data to a canonical format and then sends it to a PlotDevice for rendering. By substituting plot devices, the user can cause the same graph to be output to a window or a file using gnuplot, fig, xgraph, xdvi, ghostview, or acroread. In Figure 1 we show a utilization graph that collects data from a link in an ns simulation and then outputs to a gnuplot plot device. The gnuplot plot device translates the data and commands from the utilization graph into gnuplot commands. The gnuplot application then generates an encapsulated postscript file. Other scenarios include rendering a graph with multiple plot devices for example to generate a postscript file while simultaneously displaying the data in an X window using xgraph. Or multiple graphs may be output to the same acroread plot device so that the graphs appear in a single acroread window in the given order.

2 Installation

Presumably you already have the distribution and have installed ns-2. The graph package has been tested with ns-2.1b9a, ns-2.26, and ns-2.28. There may be incompatibilities with other ns versions. If you have not downloaded the graph package then get it from

<http://www.ecse.rpi.edu/~harrisod/graph.html>.

*I have graduated from RPI. Further maintenance and development is being done via Source Forge.
<http://www.sourceforge.net/projects/ns2graph>

¹Since the name “RPI ns-2 Graphing and Statistic Package is bulky, we use the terms “RPI graph package” and “graph package” interchangeably with the full name.

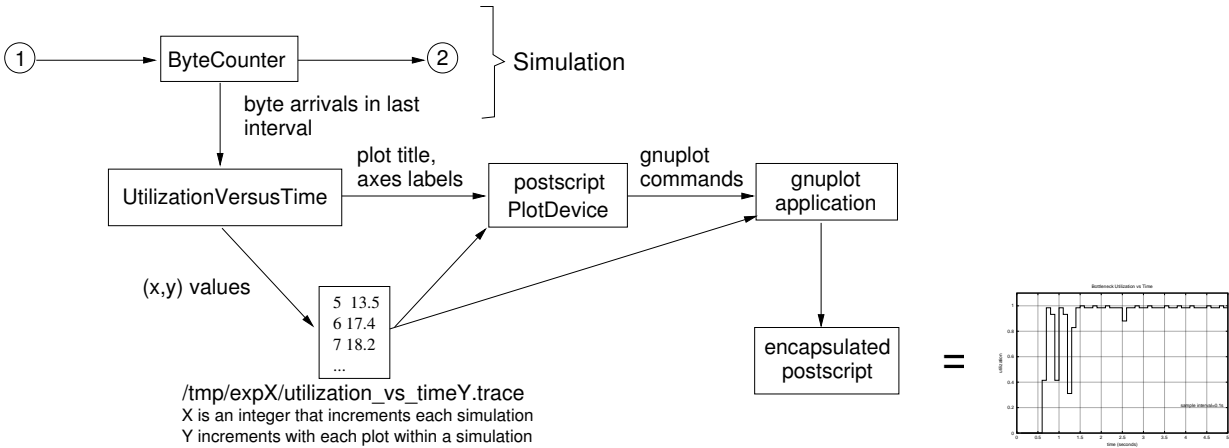


Figure 1: UtilizationVersusTime Graph outputting via postscript PlotDevice

The graph package comes with an `install` script, which is found in the root of the directory of the distribution. This `install` script will help you through the installation process. However, it is not fully automatic. You will have to edit your shell initialization script (e.g., `.bashrc`) and the NS Makefile manually. The install script will prompt you when this is necessary.

To unpack the distribution and to run the `install` script enter

```
> gunzip graph-v6.x.tar.gz
> tar -xf graph-v6.x.tar
> cd graph-v6.x
> ./install
```

where x is the minor version number of the graph package.

Should the install script fail you may use the manual install procedure given in the next section.

3 Manual Installation

The following steps are the same as those taken by `install`:

1. In order for the graph package to function, you must define two environment variables: the `NS` environment variable must point to the root of the `ns` source code tree, and the `NSVER` environment variable must contain the `ns` version number (e.g., `ns-2.1b9a` has version “2.1b9a”). If you are using `bash` then add adding the following to your `.bashrc` file

```
export NS=/home/harrisod/ns-allinone-2.28/ns-2.28
export NSVER=2.28
```

but change `/home/harrisod` to the appropriate path and `NSVER` to reflect the appropriate NS version number.

2. Make a backup of your `ns` source code tree (i.e., `$NS`) before attempting this installation.

3. If you have an earlier version of the graph package then remove it as follows

```
> cd $NS
> rm -r rpi
> rm -r tcl/rpi
```

4. Move directories from the distribution into your \$NS distribution.

```
> cp -r rpi $NS
> cp -r tcl/rpi $NS/tcl/rpi
> cd $NS
```

5. To configure the paths to the applications used by the graph package run the `configure.tcl` script in `$NS/tcl/rpi`,

```
> cd $NS/tcl/rpi
> ns $configure.tcl
```

This script will look for the following applications:

- `nam`,
- `acroread`,
- `epstopdf`,
- `ghostview`,
- `gnuplot`,
- `latex`,
- `pdflatex`,
- `xdvi`,
- and `xgraph`.

In some cases, the script will also look for alternate applications. For example in Linux, `evince` can display both postscript and pdf files, and thus can be used in place of `ghostview` and `acroread`.

`nam` comes with `ns-2`, is necessary to run network animations, and is used in some of the examples provided with this package (see the `examples` directory in the root directory of the graph distribution). The remaining applications are used in the process of creating graphs. Not all of the applications are necessary to use the graph package though some `PlotDevice` classes will not function. See Table 2 for the list of `PlotDevice` classes and the dependencies between these classes and the applications.

6. Insert the following lines in your `ns` Makefile

```

[...]
INCLUDES= \
    -I. -Irpi \
[...]
OBJ_CC= \
[...]
rpi/byte-counter.o rpi/delay-monitor.o rpi/file-tools.o \
    rpi/rate-monitor.o rpi/rpi-flowmon.o rpi/rpi-queue-monitor.o \
[...]

```

Here “INCLUDES= \” and “OBJ_CC= \” should already appear in your Makefile. “[...]” refers to an omission of lines already appearing in your Makefile.

7. Rebuild ns

```

> cd $NS
> rm gen/*
> make depend
> make

```

8. Test ns

Change to the directory \$NS/tcl/rpi/tests. Type

```
run-test-suite.sh
```

This will test the RPI Graphing and Statistics package. You should see the following output:

```

script-tools.tcl Test:  PASSED all 17 tests.
ByteCounter Test:      PASSED
DelayMonitor Test:     PASSED all 6 tests.
DelayMonitor Garbage Collection Test :  PASSED
file-tools.tcl Test:   PASSED
RPIQueueMonitor Tests: PASSED all 24 tests.
link-stats Tests:     PASSED all 59 tests.

```

After the text above, the test suite calls `graph.test/graph-test.tcl`, which generates a set of graphs and displays them using a variety of plot devices.

Carefully inspect the generated plots for correctness against the plots bearing the label “Comparison Graph.” Because of the visual nature of the graphing tools and slight differences between versions of the various graphing applications, we found that the only reliable way to test the Plot-Device classes was through visual inspection. NOTE: That your output may not look exactly the same as the provided comparison graphs. Look for differences in content rather than small differences in presentation (e.g., ignore font differences).

4 Tutorial

All of these examples assume a working knowledge of **ns**. If you have not created **ns** scripts before then consult the **ns** documentation and write a few test scripts before proceeding from here.

Currently the graph package provides the graphs shown in table 1.

Except for the **XY** class, an instance of any of these classes gathers statistics directly from an **ns** simulation and renders a plot when the instance's **plot** member function (i.e., TCL instproc) is called. The **XY** class simply allows script writers to generate a plot from an arbitrary set of (x, y) coordinates. The **XY** class represents a way for a script writer to use the **PlotDevice** classes as a generic interface to a variety of output devices.

CWndVersusTime	TCP congestion window size versus time
FlowQLenVersusTime	sampled per-flow queue contribution versus time
PointToPointDelayVersusTime	delay between two points on the network
QDelayVersusTime	sampled queuing delay versus time
QLenVersusTime	sampled queue length versus time
RateVersusTime	arrival rate versus time over constant interval length
REDQueueVersusTime	avg queue, instant queue versus time
RTTVarianceVersusTime	TCP variance in round-trip-time versus time
RTTVersusTime	TCP round-trip time versus time
Sequence	TCP sequence number versus time
SRTTVersusTime	TCP smoothed round-trip time versus time
UtilizationVersusTime	utilization versus time over constant interval length
XY	plot (x, y) coordinates from an input file

Table 1: Provided Graph classes

4.1 Example 1: Graphing link statistics

To include the graph package we preface every experiment script with the following:

```
source $env(NS)/tcl/rpi/graph.tcl
```

The graph package contains several graphs for gathering statistics about traffic passing through a link including queue length versus time, utilization versus time, rate versus time, and queue delay versus time. In this example we will show how to graph utilization over given fixed intervals versus time.

Assume we have created a topology with a known bottleneck spanning **ns** nodes **n0** and **n1**. Next we add the lines shown in figure 2 to our script. To collect statistics for our graph we instantiate a **Graph/UtilizationVersusTime** object. The arguments **\$n0** and **\$n1** tell the graph object to collect statistics for the link spanning between nodes **n0** and **n1**. The 0.1 argument tells the graph object that it should measure utilization over 0.1 second intervals. When the simulation finishes, we call **\$util_graph display** to tell the graph object to show a window containing our utilization versus time plot generated using **gnuplot**. We show the complete source code for this script in **graph-v6.x/examples/ex1.tcl**, and the generated graph in figure 3.

```

[...]
```

```

proc finish {} {
    global util_graph
    [...]
    $util_graph display
    exit 0
}

[...];# define nodes n0 and n1 and link spanning them.
# create graph of utilization vs time for link [n0,n1].
set util_graph [new Graph/UtilizationVersusTime $n0 $n1 0.1]
```

Figure 2: Utilization versus time script: ex1.tcl

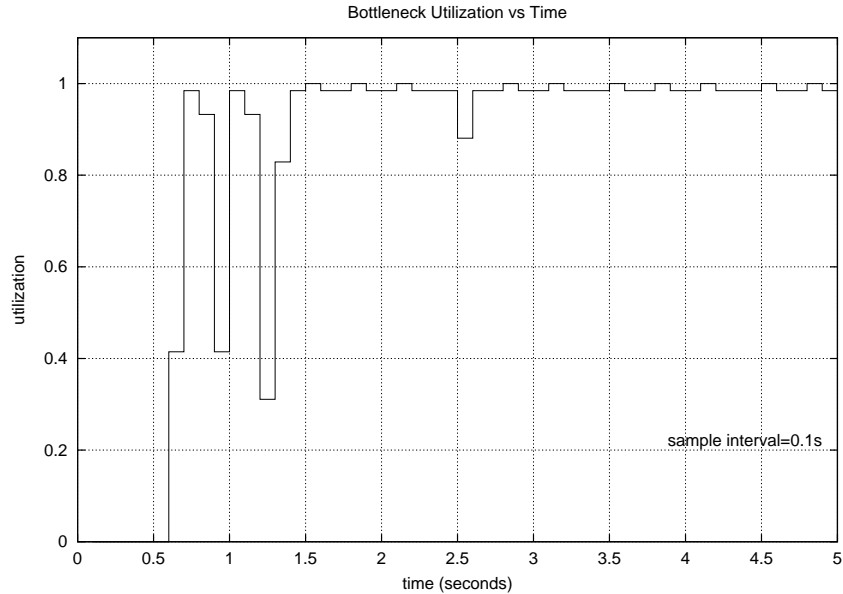


Figure 3: Graph generated by ex1.tcl

4.2 Example 2: Outputting to Various Plot Devices

The graph package can output graphs to windows or files using various applications or formats for displaying/saving graphs. We show the current set of supported “plot devices” in table 2.

In `graph-v6.2/examples/ex2.tcl` we provide an example script that uses an `xdvi` plot device to display multiple graphs in the same window in the designated order. One of the first things we do in this script is to create a plot device:

```
Graph set plot_device_ [new xdvi]
```

Here we create an instance of the `xdvi` class and assign it to the `Graph plot_device_` class member. Whenever a graph object is instantiated it by default uses the plot device defined in the

acroread	When closed, display all graphs plotted with this device using acroread. Uses postscript PlotDevice and the pdflatex application.
fig	Output to fig file via gnuplot.
ghostview	Output eps via gnuplot then display using ghostview
gnuplot	Show plot in X window using gnuplot.
gnuplot35	Show plot in X window using gnuplot 3.5 (no comments).
latex	Output to eps file using the postscript PlotDevice and run latex to create a latex file that includes the graphs.
pdf	Output to pdf file from eps using epstopdf.
postscript	Output to encapsulated postscript via gnuplot.
postscript35	Output to eps via gnuplot 3.5 (no comments).
xdvi	When closed, display all graphs plotted with this device using xdvi.
xgraph	Show plot in X window using xgraph.

Table 2: Supported plot devices

Graph `plot_device_` class member. All of the classes can be instantiated by passing a plot device object to their respective `init` instprocs. This allows you to use a different plot device for each graph object if you so wish.

We show the relevant snippets from the `ex2.tcl` script in figure 4. Not shown in figure 4 we create a topology containing a bottleneck between nodes `$n0` and `$n1`. We then define two TCP connections that pass through the bottleneck. The first of these two connections has the TCP agent object `$tcp0`. Shown in figure 4, we create three graph objects of which the first two install statistics gathering objects into the bottleneck link from `n0` to `n1`. The last graph object traces the congestion window of `tcp0`. When the simulation completes we output each of the graphs to encapsulated postscript by calling each object's `display` instproc. Next we close the plot device causing the plot device to compile a latex file that references the eps files generated by each graph object. The latex compiler outputs an dvi file which the plot device then opens using `xdvi`.

4.3 Example 3: Outputting Graphs as Files

Often we want to generate graphs as files for inclusion in a report or paper. In the previous example we showed how to output into `xdvi`. In this example we consider outputting to a `fig` file for later annotation using `xfig`. With this example we also take into consideration where generated files are placed in the file system.

We show code snippets for Example 3 in Figure 5. Example 3 differs from Example 2 only in that we use a different plot device, in this case `fig`, and we explicitly tell `util_graph` to output to the file named

```
bneck_util_vs_time
```

in the current working directory. Note that this is not the exact filename of the output file. The plot device appends an id unique to the plot device and then appends a file type extension. The actual output file's name is

```

[...]  

Graph set plot_device_ [new xdvi]  

[...]  

proc finish {} {  

    global util_graph qlen_graph cwnd0_graph  
  

    $util_graph display  

    $qlen_graph display  

    $cwnd0_graph display  
  

    [Graph set plot_device_] close  
  

    run-nam  

    exit 0  

}  

[...];# define a bottleneck between n0 and n1.  
  

set util_graph [new Graph/UtilizationVersusTime $n0 $n1 0.1]  

$util_graph set title_ "Bottleneck Utilization vs Time"  
  

set qlen_graph [new Graph/QLenVersusTime $n0 $n1]  

$qlen_graph set title_ "Bottleneck Queue Length Versus Time"  
  

set cwnd0_graph [new Graph/CWndVersusTime $tcp0]  

$cwnd0_graph set title_ "cwnd of flow 0 versus Time"  

[...]
```

Figure 4: Example 2: Display multiple graphs using xdvi

bneck_util_vs_time_plot1.fig

Note that setting `output_filename_` does not affect the location of any trace or other temporary files. All temporary files are by default written to a temporary directory located in `/tmp/exp x` where x is replaced with the smallest integer that has not already been used in the naming of another `exp` directory in `/tmp`. The user can change the default directory by setting `tmp_directory_` global variable before instantiating any Graph objects. However, this is only advisable if the user is sure that the directory used to store temporary files resides on the same machine that is executing the script, since writing across a network will not only slow down the simulation but may adversely affect other users sharing the network.

If the user does not set the `output_filename_` data member then the output file is placed in `tmp_directory_` with a default name specific to the Graph class.


```
[...]
Graph set plot_device_ [new fig]
[...]
set util_graph [new Graph/UtilizationVersusTime $n0 $n1 0.1]
$util_graph set title_ "Bottleneck Utilization vs Time"
$util_graph set output_filename_ "bneck_util_vs_time"
[...]
```

Figure 5: Example 3: Outputting Graphs to Files

4.4 Example 4: Multiple Plot Devices

Note that the caller can output the same graph object to both a file and a window by using two plot devices. In `ex4.tcl` we use `xgraph` and `postscript` to generate encapsulated postscript while displaying the graph using `xfig`. We show the relevant code snippet in Figure 6. Notice that we do not call the “display” instance procedure on the graph object. The graph’s “display” method plots using the default `PlotDevice`. Instead we call the plot device directly using its “plot” instproc, and we pass the graph to the plot device.

Note that often times there is little need for more than one plot device since plot devices create intermediate files of the desired type (e.g., `xdvi` creates encapsulated postscript files for each graph). Multiple plot devices is more useful when the desired file type is not generated as an intermediate file when displaying the graph in a window as in `ex4.tcl`.

```
[...]
proc finish {} {
    [...]
    set xgraph_plotter [new xgraph]
    set eps_plotter [new postscript]
    $xgraph_plotter plot $util_graph
    $eps_plotter plot $util_graph

    $xgraph_plotter close
    $eps_plotter close
    [...]
}
[...]
set util_graph [new Graph/UtilizationVersusTime $n0 $n1 0.1]
$util_graph set title_ "Bottleneck Utilization vs Time"
$util_graph set output_filename_ "bneck_util_vs_time"
[...]
```

Figure 6: Example 4: Output to a file and to an X window using `xgraph`

4.5 Example 5: Using overlays

With all of the graph classes in the graph package you can overlay the data from one plot on top of another. For example, in this section we show how to plot congestion window versus time for two TCP connections in the same plot. In Figure 7 we show a TCL snippet that calls `overlay` in the `finish` procedure to combine `cwnd0_graph` and `cwnd1_graph`.

```
[...]
proc finish {} {
    [...]
    $cwnd0_graph overlay $cwnd1_graph "TCP 1 " "TCP 0 "
    $cwnd0_graph display
    [...]
}
[...]
set cwnd0_graph [new Graph/CWndVersusTime $tcp0]
set cwnd1_graph [new Graph/CWndVersusTime $tcp1]
$cwnd0_graph set title_ "cwnd of flows 0 and 1 versus Time"
[...]
```

Figure 7: Example 5: overlay cwnd for connection 1 on cwnd graph for connection 0.

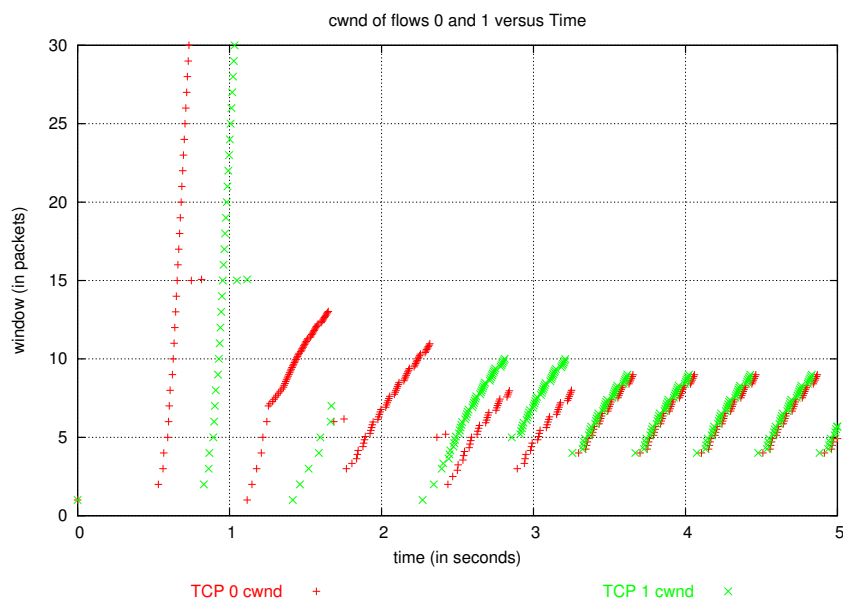


Figure 8: Graph generated by `ex5.tcl`

The act of overlaying plot x onto plot y by calling `$y overlay $x` causes x 's `prepare` method to be called, but it does not otherwise affect x . However, y copies x 's data set objects. y keeps its own data sets. However, y and x share x 's data files for the data sets obtained from x . y retains

its basic appearance. Applying an overlay does not change y 's title, axis labels, caption, comment, over horizontal lines. If the x or y axis ranges have been defined then they also do not change. However, if the x or y axis ranges have not been explicitly set then the underlying plot device may modify the axes ranges to display all data sets. For example, gnuplot and xgraph scale x and y axes automatically when the ranges are not explicitly provided.

The `overlay` procedure takes three arguments. The first argument specifies the graph that will be overlaid. The second argument specifies a prefix that is prepended to all data set names (i.e., the names that appear in the key) for the data sets obtained from the graph specified in the first argument. In Figure 7, “TCP 1 “ is prepended to connection 1’s cwnd label in the key. The third argument specifies a prefix that is prepended to all of the data sets in the graph whose `overlay` procedure is being called. In Figure 7, “TCP 0 “ is prepended to connection 0’s cwnd label in the key. The last two arguments are optional.

4.6 Example 6: Bottleneck link statistics

We now move on to demonstrating how to collect statistics using the statistics functions included with the graph package. The objective of these statistics gathering tools is to simplify statistics gathering over the standard methods provided by `ns`.

To include the link statistics gathering functions do the following:

```
source $env(NS)/tcl/rpi/link-stats.tcl
```

To begin gathering statistics on a link spanning from nodes `n0` to `n1`, do the following:

```
set stats [new LinkStats $n0 $n1]
```

We show code snippets from `examples/ex5.tcl` in figure 9. The calls to `get-utilization`, `get-mean-queue-delay`, and `get-packet-arrivals` returns the link statistics gathered from the moment the link-stats object was instantiated. In this case, that means reporting statistics gathered from the beginning of the simulation. There are many more statistics we could report by simply calling any of the instance procedures in table 4.

If you wish to gather statistics starting from some time into the simulation then simply instantiate the `LinkStats` object at that time. If you want to gather statistics over consecutive time intervals then you can reset the link statistics at the end of each time interval by calling the `LinkStats reset` instance procedure. If you want to collect a link’s statistics for overlapping intervals then create two `LinkStats` objects each at a time when you want to begin collecting statistics then retrieve the statistics at the end of their respective intervals. Installing more than one `LinkStats` object on a link is particularly useful when you want to record statistics for the whole simulation as well as for the steady-state (i.e., skipping some time to eliminate initial transients).

We provide the code for this complete example in `graph-v6.2/examples/ex6.tcl`.

4.7 Example 7: TCP statistics

The most direct way to obtain TCP statistics is to simply query the bound variables defined for the Agent/TCP class. For example, if we define a TCP agent `$tcp0` then we can do the following without extending `ns-2`:

```

[...]
proc finish {} {
    global n0 n1 tcp0 tcp1 stats

    # output link statistics
    puts "Bottleneck statistics: "
    puts "Utilization: [$stats get-utilization]"
    puts "Number of drops: [$stats get-packet-drops]"
    puts "Mean queue delay [$stats get-mean-queue-delay]"
    puts "Number of arrivals [$stats get-packet-arrivals]"
    [...]
}
[...] ;# create topology w/ bottleneck between n0 and n1.

set stats [new LinkStats $n0 $n1]
[...]
```

Figure 9: Example 6: Gathering link statistics

<code>init <i>n0 n1 [qmon]</i></code>	Instantiates a LinkStats object spanning between <i>n0</i> and <i>n1</i> . You must have previously created these two nodes and a link spanning these nodes. More than one LinkStats object can span between <i>n0</i> and <i>n1</i> . By default the LinkStats object uses an instance of the RPIQueueMonitor class to collect statistics. You can install your own queue monitor using the optional third argument.
<code>reset</code>	resets all statistics to the initial state.

Table 3: Special instance procedures for LinkStats

```

puts "Number of data packets transmitted by flow 0: \
    [$tcp0 set ndatapack_]"
puts "Number of data packets transmitted by flow 1: \
    [$tcp1 set ndatapack_]"
puts "Retransmission timeouts for flow 0: \
    [$tcp0 set nrexmit_]"
puts "Retransmission timeout for flow 1: \
    [$tcp1 set nrexmit_]"
```

When using our TCP statistics gathering functions it is necessary to include the `tcp-stats.tcl` package. This loads the definitions shown in tables 5 and 6:

```
source $env(NS)/tcl/rpi/tcp-stats.tcl
```

Then initialize various counters by calling `init-stats`. `init-stats` can also be called to simply reset the the statistics for both our extensions and the statistics gathered by `ns`.

```
[...]
$tcp init-stats
[...]
```

At some later time, such as when `finish` is called, you can output tcp-statistics as follows:

```
puts "Goodput Variance: \
    [get-goodput-stddev "$tcp0 $tcp1"]"
puts "Total TCP packets: \
    [get-total-data-packets "$tcp0 $tcp1"]"
```

In the example above, `get-goodput-stddev` and `get-total-data-packets` operate on a list of tcp agents. In this case returning the standard deviation in goodput and the sum of the data packets sent respectively across the passed list of agents. The above examples can be found in `examples/ex7.tcl`.

<code>get-utilization</code>	Returns the link utilization as the number of bytes departing the queue as the number of bytes that could have departed the queue.
<code>get-packet-utilization</code> <i>avgpkt sz</i>	Returns the link utilization as the number of packets that departed the bottleneck link times the average packet size over the number of average sized packets that could have departed the link. The argument specifies the average packet size.
<code>get-throughput</code>	Returns the number of byte arrivals at the tail of the queue * 8 over time.
<code>get-packet-arrivals</code>	Returns the number of packet arrivals at the tail of the queue.
<code>get-byte-arrivals</code>	Returns the number of byte arrivals at the tail of the queue.
<code>get-packet-drops</code>	Returns the number of packets dropped.
<code>get-byte-drops</code>	Returns the number of bytes dropped.
<code>get-packet-departures</code>	Returns the number of packets that departed the queue.
<code>get-byte-departures</code>	Returns the number of bytes that departed the queue.
<code>get-mean-queue-delay</code>	Returns the mean delay experienced by packets transiting the queue.
<code>get-queue-delay-variance</code>	Returns the variance in delay experienced by packets transiting the queue.
<code>get-queue-delay-stddev</code>	Returns the standard deviation in delay experienced by packets transiting the queue.
<code>get-mean-packet-queue-length</code>	Returns the mean packet queue length. This is not the same as the mean queue length seen by an arriving packet which would not take into account idle times in the mean.
<code>get-mean-byte-queue-length</code>	Same as <code>get-mean-packet-queue-length</code> except the result is returned in units of bytes.
<code>get-max-packet-queue-length</code>	Returns the maximum queue length in packets.
<code>get-max-byte-queue-length</code>	Returns the maximum queue length in bytes.
<code>get-min-packet-queue-length</code>	Returns the minimum queue length in bytes. Typically this would be zero except when measured over some small interval.
<code>get-min-byte-queue-length</code>	Same as <code>get-min-byte-queue-length</code> except the result is returned in units of bytes.

Table 4: LinkStats statistics instance procedures

<code>get-useful-packets</code>	Returns the number of packets transmitted containing new data.
<code>get-useful-bytes</code>	Returns the number of bytes transmitted containing new data. The number of useful bytes in each useful packet is the packet size minus the TCP and IP headers. The header size is determined by <code>tcip_base_hdr_size_</code> defined in <code>\$NS/tcl/lib/ns-default.tcl</code> . In ns-2.1b5, <code>tcip_base_hdr_size_</code> is set to 40.
<code>get-goodput-bps</code>	Returns the rate of useful bits transmitted by the network. Uses <code>get-useful-bytes</code> to determine the total number of useful bits. Packets that have been lost but not yet detected by the source are counted as useful because our measure is based on state maintained by the TCP source. In a long simulation this should have negligible impact on the goodput measure.
<code>get-goodput</code>	Returns goodput in bps over the bottleneck capacity.
<code>get-throughput-bps</code>	Returns the throughput in bps. This counts each retransmission as a separate packet. This differs from <code>get-useful-bytes</code> and <code>get-useful-packets</code> which only count the packets sent containing new data.

Table 5: TCP statistics gathering instance procs

<code>get-mean-goodput <i>tcplist</i></code>	Calculates mean goodput in bps across the passed TCP agents (i.e., connections).
<code>get-goodput-variance <i>tcplist</i></code>	Calculates variance in goodput (bps) across the passed TCP agents (i.e., connections).
<code>get-goodput-stddev <i>tcplist</i></code>	Calculates the standard deviation in goodput in bps across the passed TCP agents (i.e., connections).
<code>get-goodput-cov <i>tcplist</i></code>	Calculates the Coefficient of Variation (C.O.V.) across the passed TCP agents. C.O.V. is standard deviation over the mean.
<code>get-total-data-packets <i>tcplist</i></code>	Calculates the sum of the packets sent across the passed TCP agents (i.e., connections).
<code>get-total-retransmitted-packets <i>tcplist</i></code>	Calculates the sum of the packet retransmissions across the passed TCP agents (i.e., connections).
<code>get-total-retransmission-timeouts <i>tcplist</i></code>	Calculates the sum of the retransmission timeouts across the passed TCP agents (i.e., connections).

Table 6: TCP statistics functions operating on lists

4.8 Example 8: Customizing graph appearance and layout

In prior sections we mostly relied on the default behavior of the Graph and PlotDevice objects. In this section we explore some of the options available for affecting the appearance and arrangement of graphs.

First we consider Graph options. In particular we consider captions, comments, axes labels, adding horizontal lines and inserting latex in the output.

Title To set the title that appears in the plot set the `title_` data member of the graph object. For example,

```
$qlen_graph set caption_ "Bottleneck Queue Length Versus Time"
```

Axis Labels To set an axis label, set the `ylabel_` or `xlabel_` data members of the graph object. For example,

```
$qlen_graph set xlabel_ "Simulation Time (in seconds)"
```

Captions To incorporate captions, define the `caption_` data member of the graph object like follows:

```
$qlen_graph set caption_ "Two TCP connections traverse this bottleneck."
```

Horizontal Lines You can add horizontal lines to any graph by calling the `add-hline` instproc. For example,

```
$qlen_graph add-hline 10 "Buffer size"
```

This adds a horizontal line at 10 and adds a entry in the key labelled “Buffer size.” This package currently provides no method for creating verticle lines.

Comments You can add a comment to any graph at a position specified using x,y coordinates that each range from 0 to 1. Note that the coordinates are not specified in data units. (0,0) places the comment near the origin, and (1,1) places the comment in the upper right-hand corner. For example, in our queue length graph we can add a comment stating that the queue length is instantaneous as opposed to averaged and place this comment in the upper left-hand corner of the plot by doing the following:

```
$qlen_graph set comment_ "Instantaneous queue length"
$qlen_graph set xcomment_ .3
$qlen_graph set ycomment_ .65
```


Axes Ranges We can control the range of values on each axes in order to zoom in on a region or to improve appearance by setting the `xhigh_`, `xlow_`, `yhigh_`, and/or `ylow_` data members of the graph object. If for example, we want to zoom in on the first two seconds of the simulation and we want to extend the upper range of the y-axis so that the buffer size horizontal line is not coincident with the top border of the graph, we add the following to our script:

```
$qlen_graph set xhigh_ 2
$qlen_graph set yhigh_ 11
```

We add all of these options together in example 8 as shown in Figure 10. We show the unadorned graph in Figure 11(a), and with all of the appearance modifications in Figure 11(b).

```
Graph set plot_device_ [new acroread]
[...]
proc finish {} {
  global qlen_graph
  [...]
  $qlen_graph add-hline 10 "Buffer size"
  $qlen_graph set title_ \
    "Bottleneck Queue Length Versus Time"
  $qlen_graph set ylabel_ \
    "Queue Length (in packets)"
  $qlen_graph set xlabel_ "Simulation Time (in seconds)"
  $qlen_graph set caption_ \
    "Two TCP connections traverse this bottleneck."
  $qlen_graph set comment_ \
    "Instantaneous queue length"
  $qlen_graph set xcomment_ .3
  $qlen_graph set ycomment_ .65
  $qlen_graph set xhigh_ 2
  $qlen_graph set yhigh_ 11
  $qlen_graph display
  [...]
}
[...] ;# create topology w/ bottleneck between n0 and n1.
```

Figure 10: Example 8: Customizing graph appearance.

Assume that we want to juxtapose plots so that more than one appears on each line (i.e., more than one per row). We do this by telling the latex class to generate appropriate latex to fit the desired number of graphs in each row of output. For example, if we want 2 graphs per row we do the following:

```
latex set n_plots_per_row_ 2
Graph set plot_device_ [new acroread]
```

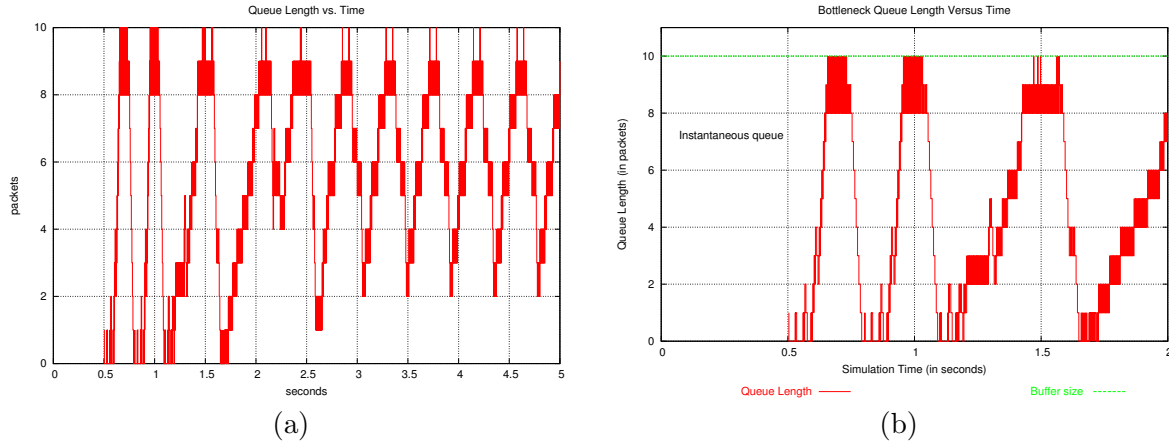


Figure 11: (a) shows the graph with default appearance. (b) shows the graph with modified appearance using the options described in Section 4.8.

Above, we set the `latex` class's `n_plots_per_row_` class variable. By setting this value before creating the `acroread` plot device, we ensure that any `latex` object instantiated by `acroread` as a low-level plotter will set its object data member `n_plots_per_row_` equal to its class data member of the same name.

If we want to insert some text anywhere in the output latex, we call the plot device's `output-latex` instproc. The text will appear in order with any plots meaning if we want the text to appear before any plots then we must call `output-latex` before calling a graph object's `display` instproc to generate the first plot. For example, if we want to add a title to the output latex file we could do the following:

```
[...]
set pd [new acroread]
Graph set plot_device_ $pd
$pd output-latex \
  {\title{Example 8 Simulation Results}\maketitle}
[...]
```

If we want to output some descriptive text in between plots we do the following:

```
$qlen_graph plot
[...]
$pd output-latex "We now want to generate another plot."
[...]
$qlen_graph plot
```

Unfortunately the `latex` plot device class is not particularly smart. It will allow you to output latex at any point even if it screws up other formatting. For example, if you are outputting more than one plot per row, be careful not to insert text in the middle of a row. For example, if `n_plots_per_row_` is 2 then only call `output-latex` after generating an even number of plots.

Note that not all of these options are supported by all plot devices. When an option is unsupported, the plot device simply ignores the option.

4.9 Example 9: XY Plots

If all existing graph types do not satisfy your needs, and you just need to graph something right now as opposed to going through the overhead of extending the graph package, we suggest using the XY graph class. The XY graph class takes as input a text file containing two white-space delimited columns. The left column contains X values and the right column contains Y values.

For example, when we create an XY graph

```
set xy_graph [new Graph/XY "ex9_thruput_vs_rtpd.txt"]
```

that XY graph expects (x, y) coordinates to be output into a file named `ex9_thruput_vs_rtpd.txt` in a format like

```
16 18530.909090909092
24 93032.727272727279
32 149381.81818181818
...
```

where 16, 24, and 32 are x -values and the right column provides the y -values.

In `ex9.tcl`, we show how to use Graph/XY to graph something that is not provided by a graph within the RPI graph package. In this particular example we demonstrate the effect of heterogeneous round-trip propagation delay on TCP. It is well-known that TCP throughput is roughly inversely proportional to round trip time.

To demonstrate the bias against long round-trip times we first create a dumbbell topology with N paths each with a different round-trip propagation delay.

```
for { set i 0 } { $i < $N } { incr i } {
    set s($i) [$ns node]
    set d($i) [$ns node]
}
set b0 [$ns node]
set b1 [$ns node]

# create links.
for { set i 0 } { $i < $N } { incr i } {
    $ns duplex-link $s($i) $b0 10M [expr 4 * $i]ms DropTail
    $ns duplex-link $b1 $d($i) 10M 4ms DropTail
}
```

We create one TCP/Reno connection on each path through the dumbbell.

```
for { set i 0 } { $i < $N } { incr i } {
    create-ftp-over-reno $s($i) $d($i) [expr 0.5 * $i] $i 0
}
```

We then install link statistics gathering components on each of the bottleneck-destination links as follows:

```
for { set i 0 } { $i < $N } { incr i } {
    set lstats($i) [new LinkStats $b1 $d($i)]
}
```

When the simulation completes after sufficient time (30 seconds) for the simulation to reach steady-state, we read the statistics from the LinkStats objects and output (*delay, throughput*) values to the file named

```
ex9_thruput_vs_rtpd.txt
```

by doing

```
for { set i 0 } { $i < $N } { incr i } {
    set thruput [$lstats($i) get-throughput]
    puts $fp "[expr ($i * 4 + 4 + 4) * 2] $thruput"
}
```

We then generate the graph just as we would with any other graph, we call the graph's `display` procedure,

```
$xy_graph display
```

A slightly more complete summary of this script is provided in Figure 12. The resulting graph is shown in Figure 13. The graph does not show a strictly decreasing throughput with increasing round-trip propagation delay, but this could be due such factors as synchronization effects because we are using a drop tail queue. The reader could remove synchronization effects by placing a random dropper like RED at the bottleneck. We leave this as an exercise.

5 Performance Issues

When one is running large or long simulations, there are various considerations to take into account with respect to the tools used for gathering statistics or generating statistics. Our tools have been defined to introduce reasonable overhead, though in some cases simulation performance constraints may drive one to design task-specific tools. We outline some of the performance considerations in this section.

5.1 When to Use NAM, Graph, and LinkStats

`nam` tracing is best used to generate animations. For any simulation containing more than a handful of nodes, containing bottlenecks with large bandwidths (> 10 Mbps), and/or running more than a couple of seconds, `nam` tracing will probably be too slow. `nam` is slow because in order to generate network animations `ns` outputs to a trace file every time a packet enters, departs, or is dropped from any queue in the simulated network. Simply turning `nam` tracing off often reduces simulation run times by an order of magnitude. In particular, if you are parsing `nam` traces to generate statistics then perhaps you should consider using some other tool.

Graph objects sit on a single link or deal with a single TCP agent. Unless the user wants to generate graphs for every link in the simulated network, installing Graph objects will probably

```

[...]
proc finish {} {
    global xy_graph lstats N

    set fp [open "ex9_thruput_vs_rtpd.txt" "w"]
    for { set i 0 } { $i < $N } { incr i } {
        set thruput [$lstats($i) get-throughput]
        puts $fp "[expr ($i * 4 + 4 + 4) * 2] $thruput"
    }
    close $fp

    $xy_graph display

    [Graph set plot_device_] close
    exit 0
}

set ns [new Simulator]
set xy_graph [new Graph/XY "ex9_thruput_vs_rtpd.txt"]
$xy_graph set title_ \
    "Connection Throughput versus Round-trip Propagation Delay"
$xy_graph set xlabel_ \
    "Round-trip Propagation Delay (in ms)"
$xy_graph set ylabel_ \
    "Throughput (in bps)"
[...]
```

Figure 12: Example 9: Generate an XY plot to show throughput versus RTPD.

yield significantly less overhead in terms of simulation run time and disk space consumption than generating graphs by post-processing `nam` traces. Of course, the `Graph` objects are best used for generating graphs, any other use of the `Graph` classes or their respective intermediate files is suspect. The `Graph` objects generate trace files (though substantially smaller than `nam` traces), and useful statistics might be derived from these traces. However, if you only want to determine a single statistic across the length of a simulation then the functions found in `link-stats.tcl` or `tcp-stats.tcl` are probably better for the job.

The `link-stats.tcl` and `tcp-stats.tcl` functions generate far less overhead than either `nam` or the `Graph` classes, simply because *none* of the `link-stats.tcl` or `tcp-stats.tcl` generate trace or other intermediate files. Instead the `LinkStats` object installs counters in the link being monitored. Incrementing a counter introduces far less overhead than outputting to a file. When the user later requests the statistic, the corresponding components are then queried. However, the `LinkStats` object installs objects for measuring a large array of link statistics. Sometimes this means installing unnecessary counters or similar objects into the links being monitored. For example, the `LinkStats` object installs objects at the tail and head of the link's queue. Placing an

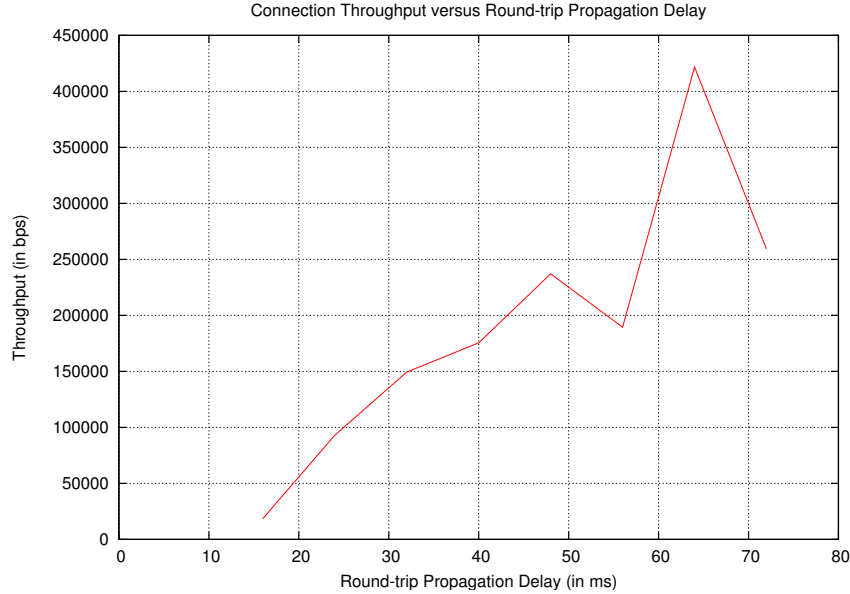


Figure 13: XY Graph generated by ex9.tcl showing the bias round-trip propagation delay has on TCP throughput.

object at the queue's head is unnecessary if the script writer only wants to know the number of bytes that arrived at the head of the queue. Usually the few extra objects represents negligible overhead. However, if the script writer has particularly tight performance constraints then the script writer will have to write his or her own objects designed specifically to gather the desired statistics.

6 Other Performance Issues with Graphs

There are two other important performance issues with regard to the graph classes: 1) where to output files, and 2) using averaging intervals.

The cost of outputting to one file is not necessarily the same as the cost of outputting to another file. The disparity is particularly large between local and remotely mounted files. A trace file should probably *never* be placed on a volume mounted across a network, because each write may generate a packet. Consider for `nam` that this would mean generating a packet on the real network for every packet arrival, departure, or drop in the simulated network. A single NAM simulation can easily swamp a file server while simultaneously drastically slowing down simulation run times.

Instead, trace and other intermediate files are usually output to a temporary directory. In UNIX, `/tmp` resides on a workstation's local hard drive; therefore, that is why the Graph package by default places all intermediate as well as output files in `/tmp` (see Example 3).

Now this brings us to the second important performance issue: averaging. Several of the Graph classes allow the user to specify whether to output instantaneous or time averaged values. For example, `QLenVersusTime`'s third argument to its `init` instproc (i.e., its constructor) is `sample_interval`. By default, the sample interval is set to -1 denoting that `QLenVersusTime` should record the queue length at every arrival and departure of a packet. This results in particularly large generated files and slow run times, but it also allows the user to easily see transient

queuing behavior. If the user wishes to improve simulation run times or reduce the size of trace files, the user can set `sample_interval` to a positive value denoting the time interval over which the queue length is averaged. At the end of each interval, the average is output to the graph's trace file. Larger averaging intervals result in shorter run times at the expense of generating graphs with smoother output (i.e., less transient behavior is revealed).

7 LinkStats and TCP Stats and Long Run Times

When simulations are run for long times or when bottleneck bandwidths are high, it is quite possible that certain signed long integers (32-bit) will overrun causing erroneous results from LinkStats objects. To fix this problem the user should perform some worst-case analysis of the number of bytes or packets that will pass through a link before running a simulation. We specifically mention “bytes” because byte counters tend to increment much faster than any other counters in a simulation. To avoid overruns, the simulation script must periodically reset the corresponding counters by calling LinkStats' `reset` member function (i.e., TCL `instproc`) or TCP's `init-stats` member function depending on the statistic that is in danger of an overrun. The simulation script can then either store the intermediate values of the statistic for later post-processing or only rely on the final value.

8 Thanks

Special thanks to Dr. Yong Xia for providing some of the code and beta-testing. Also thanks to Professor Shivkumar Kalyanaraman for supporting this work and for the patience of all of those RPI networking students that were inflicted with this package when it was in its early versions.